

SUPERSTAFF

AI digital employees that never sleep

Comprehensive Launch Brief

Project status, gap analysis, and stakeholder briefs
for Investors, Admins, and Users

Component	Status	Notes
Marketing landing page	READY	11 sections, fully responsive
Lead capture (DB + form)	READY	Validated, RLS enabled
Scripted chat demo	READY	10 keyword rules + fallback
Real AI chat engine	MISSING	Currently scripted only
Customer dashboard	MISSING	No admin/customer portal
Authentication	MISSING	No login system
Knowledge ingestion	MISSING	No doc upload pipeline
Embeddable widget	MISSING	No script tag to deploy
Billing / payments	MISSING	Pricing shown, not charged
Admin lead inbox	MISSING	Leads stored, no UI to view

Executive Summary

Superstaff is positioned as an AI-powered digital employee that businesses can deploy on their website 24/7 to answer customer questions, capture leads, and assist staff. What exists today is a **polished, conversion-optimized marketing site** with a working lead-capture backend and a scripted chat demo. What it is **not yet** is a deliverable product. To launch, we need to build the actual AI engine, an embeddable customer-facing widget, an admin/customer dashboard, authentication, and billing.

What we have today

- 11-section landing page with violet/navy futuristic design system
- Working lead capture (Name, Email, Company, Website) saving to a Supabase **leads** table with RLS
- Animated chat mock-up in hero + interactive scripted live demo (10 keyword rules)
- Pricing tiers (Starter \$49 / Growth \$149 / Pro \$399) wired to lead form (no checkout)
- Mobile-first responsive, SEO meta tags, semantic HTML, scroll animations

What blocks a real launch

- **No real AI:** the chat is hard-coded keyword matching — not LLM-backed
- **No way to onboard a customer:** no signup, no dashboard, no embed widget
- **No knowledge ingestion:** nothing to upload PDFs / URLs / FAQs into
- **No billing:** tiers are marketing copy only — no Stripe/Paddle integration
- **No admin view:** captured leads sit in DB with no interface to read them

Recommendation

Treat the current build as the **marketing layer (Phase 0)**. Launch in three phases: (1) ship the marketing site + admin lead inbox now to start collecting demand; (2) build the embeddable AI widget + customer dashboard for a closed beta in 6–8 weeks; (3) add billing + self-serve onboarding for public launch in 10–12 weeks.

INVESTOR BRIEF

The Opportunity

Every business with a website is paying humans to answer the same questions over and over — or losing leads because no one is online. Superstaff replaces a \$30–60k/year customer-facing hire with a \$49–399/month AI employee that works 24/7, in 50+ languages, trained on the company's own knowledge.

Product Vision

- **Hire, don't configure.** Position the product as an employee, not a chatbot.
- **White-labeled intelligence.** Powered by IAM.ng under the hood, branded fully as Superstaff.
- **Three jobs in one:** customer support, sales/lead capture, internal staff assistant.
- **Outcome-priced:** framed as 'less than one salary' — not per-message billing.

Traction & Current Asset

A production-grade marketing site is live with a working lead-capture pipeline. Every primary CTA opens a validated form that writes to a secure database (RLS-protected). This is sufficient to begin demand generation, capture a waitlist, and validate willingness to pay before heavy product investment.

Business Model

Tier	Price	Target Customer	Margin Profile
Starter	\$49 / mo	Solo founders, small sites	High (low usage)
Growth	\$149 / mo	SMBs, growing teams	Strong (sweet spot)
Pro	\$399 / mo	Mid-market, custom branding	Premium (managed)

Competitive Positioning

Generic AI (ChatGPT) is great at general questions but doesn't know your business, doesn't live on your website, doesn't capture leads, and isn't branded as your company. Purpose-built competitors (Intercom Fin, Drift, Ada) are enterprise-priced and complex to deploy. Superstaff wins on **5-minute setup**, **SMB-friendly pricing**, and the **'hire an employee' framing**.

Capital Required to Reach Revenue

- **Phase 1 (4 weeks):** AI chat engine via Lovable AI Gateway + admin lead inbox — minimal spend

- **Phase 2 (6–8 weeks):** Embeddable widget, customer dashboard, knowledge ingestion (RAG)
- **Phase 3 (10–12 weeks):** Stripe billing, multi-tenant routing, usage metering, analytics
- **Engineering:** 1 senior full-stack + 1 AI engineer for ~3 months to MVP-with-revenue

Risks & Mitigations

- **Hallucination risk** — mitigated via strict RAG on customer's own docs + escalation to humans
- **Commoditization by OpenAI** — mitigated by owning the deployment layer + customer data + workflow
- **White-label dependency on IAM.ng** — mitigated by abstracting the model layer to swap providers

ADMIN BRIEF

For the team operating Superstaff

What is currently operable

- **Marketing site:** live, single page (/) with 11 sections — Navbar, Hero, Social Proof, Use Cases, How It Works, Live Demo, Comparison, Testimonials, Pricing, FAQ, Final CTA, Footer.
- **Lead capture:** every primary CTA opens a dialog (Name, Work Email, Company, Website). Validated with Zod, written to **public.leads** in Supabase.
- **Lead schema:** id, name, email, company, website, plan, source, created_at. RLS allows anonymous inserts; reads are blocked.
- **Source attribution:** the **source** column captures which CTA was clicked (e.g. *hero-primary*, *pricing-growth*, *final-cta*) and **plan** captures the tier interested in.

What admins cannot do yet

- **View leads:** there is no admin UI. RLS blocks SELECT for everyone. Reads must currently happen via the backend dashboard.
- **Export leads:** no CSV download.
- **Edit FAQ / Pricing / Testimonials:** all content is hard-coded in components.
- **Deploy a real chat:** the live demo is scripted — there is no LLM call.
- **Onboard a customer:** no signup, no tenant model, no widget to ship.

Immediate operational tasks

#	Task	Owner	Effort
1	Build internal /admin route gated by auth + role	Engineering	1 week
2	Wire SELECT policy for 'admin' role using has_role()	Engineering	1 day
3	Lead inbox table view: filter by source, plan, date; CSV export	Engineering	3 days
4	Email notification on new lead (edge function + Resend)	Engineering	1 day
5	Replace placeholder testimonials with real customer quotes	Marketing	ongoing
6	Replace logo placeholders in Social Proof strip	Marketing	1 day
7	Add Google Analytics / Plausible for conversion tracking	Marketing	0.5 day

8	Configure custom domain + publish	Ops	0.5 day
---	-----------------------------------	-----	---------

Security & data posture

- Anonymous keys are public (expected); private keys must stay server-side.
- **leads** table: anonymous INSERT only — good. SELECT/UPDATE/DELETE are all denied.
- Add a user_roles table + has_role() function before adding any admin-only feature.
- Auto-confirm email signups should remain OFF in production.
- Before launch: enable Supabase email rate-limiting + CAPTCHA on the lead form.

Deployment checklist

- Verify SEO: title, description, OG tags, JSON-LD already present in index.html
- Confirm /robots.txt allows indexing
- Add a privacy policy + terms page (currently missing) — required before paid ads
- Set up custom domain in Lovable Cloud
- Connect a custom email domain for lead notification emails
- Smoke test: submit a lead from each CTA, confirm row appears in DB with correct source

USER BRIEF

For the businesses 'hiring' a Superstaff

What you're getting (the promise)

A trained AI employee that lives on your website 24/7. It answers customer questions in your voice, qualifies and captures leads, escalates the hard stuff to your human team with full context, and costs less than a single salary.

Three jobs your Superstaff does

Customer Support	Sales Assistant	Internal Assistant
Answers FAQs, refunds, shipping, hours — instantly, in 50+ languages.	Qualifies visitors, captures name + email + intent, hands off to your CRM.	Helps your team find SOPs, policies, product specs from your knowledge base.

How onboarding will work (when fully launched)

- **1. Train** — Upload PDFs, paste URLs, sync your help center. Setup ~5 minutes.
- **2. Deploy** — Paste one script tag onto your site. Your Superstaff goes live.
- **3. Scale** — Watch conversations, leads, and resolutions in your dashboard. Tweak its tone, knowledge, and escalation rules.

What you can do today

Right now, the public site lets you **request access**. Click any 'Hire your Superstaff' button, fill in your name, work email, company, and website, and the team will reach out within 24 hours. You can also try the scripted live demo on the homepage to get a feel for the interaction style.

What to expect after submitting a lead

- Confirmation toast on screen ('You're in. We'll be in touch within 24 hours.')
- Personal follow-up from the Superstaff team within one business day
- A scoping call to understand your business and what knowledge to ingest
- A guided setup of your first Superstaff (currently white-glove; self-serve coming soon)

Pricing at a glance

Plan	Price	Best for
Starter	\$49 / mo	Solo founders, small sites, up to 500 conversations/mo
Growth	\$149 / mo	Growing teams — unlimited conversations + lead capture (most popular)
Pro	\$399 / mo	Custom branding, dedicated success manager, full intelligence

All plans: monthly, cancel anytime, no setup fees, no credit card to start.

Roadmap to Real Launch

Phase 0 — Now (DONE)

- Marketing site live with lead capture
- Brand identity + design system established
- Demand-generation funnel ready

Phase 1 — Soft launch (2–4 weeks)

- Admin dashboard at /admin with auth + role-based access
- Lead inbox: list, filter, export CSV, email notifications on new lead
- Replace scripted demo with real LLM call (Lovable AI Gateway, no API key needed) — gemini-2.5-flash for cost
- Add privacy policy + terms pages
- Publish to custom domain, start paid acquisition

Phase 2 — Closed beta (6–8 weeks)

- Customer signup + multi-tenant model (organization per customer)
- Customer dashboard: upload knowledge (PDFs, URLs), view conversations, view captured leads
- Knowledge ingestion pipeline: chunking + embeddings + vector search (RAG)
- Embeddable widget: single <script> tag deploys a chat bubble onto any site
- Conversation logs + escalation-to-human flow (email/Slack)
- Onboard 10–20 design partners at discounted pricing

Phase 3 — Public launch (10–12 weeks)

- Stripe billing on the three tiers + usage metering
- Self-serve signup → train → deploy in 5 minutes (no human in the loop)
- Analytics dashboard: conversations, deflection rate, leads captured, top topics
- CRM integrations (HubSpot, Salesforce) + Zapier/webhooks
- Multi-language UI in widget, custom branding for Pro tier
- Public launch on Product Hunt + paid acquisition at scale

Bottom Line

Today: a beautifully presented promise with a working demand-capture funnel. To launch as a real product, the next 3 months of engineering should focus on the AI engine, the embeddable widget, the customer

dashboard, and billing — in that order. Start collecting leads now to validate pricing and ICP while building.